The paper [1] entitled "Evaluating the Suitability of Serverless Computing for Data Processing" explores the use of serverless computing for data processing through a system called Lambada. The authors evaluate the suitability of serverless computing for interactive data analytics on cold data and provide insights into the challenges faced and their solutions.

The key contribution of the paper is the proposal of a two-level invocation approach for invoking serverless workers in a batch process to reduce latency. The authors also present a family of exchange operators for serverless workers as a building block for data processing. The two optimizations proposed by the authors address the limitations of the basic exchange algorithm.

The primary result of the paper is the evaluation of Lambada in comparison to two well-established Query-as-a-Service systems, Amazon Athena and Google BigQuery. The results show that Lambada is competitive with Athena and BigQuery and well suited for interactive data analytics on cold data.

One limitation of the paper is the usage-based billing model in cloud computing, which the authors highlight as a challenge. The authors did not provide a detailed discussion on this limitation and its impact on the suitability of serverless computing for data processing.

In conclusion, the paper provides a useful idea for data processing in a serverless worker environment and contributes to the ongoing discussion on the suitability of serverless computing for data processing.

References:

[1] Müller, R. Marroquín, and G. Alonso. 2020. Lambada: Interactive Data Analytics on Cold Data Using Serverless Cloud Infrastructure. In ACM SIGMOD. 115–130.